

Bitonic Sort: Parallel Analysis

Toby Isaac

CSE 6230, Fall 2017

1 Pseudocode

Algorithm 1: `sortBitonic(r, P, n, A, d, b)`

Data: r is the process's rank;

P is the number of processes (power of 2);

n is the number of local keys (same on all processes);

A is the array of keys;

d is `True` if sorting forward, `False` if sorting backward;

b is `True` if A is bitonic on input.

```
1 if  $P = 1$  then
2   if  $b$  then
3     /* local sort when input is bitonic,  $O(n \log_2 n)$  */
4     sortLocalBitonic( $n, A, d$ )
5   else
6     /* local sort with no special condition,  $O(n \log_2 n)$  */
7     /* (but we expect slower than sortLocalBitonic()) */
8     sortLocal( $n, A, d$ )
9 else
10  if not  $b$  then
11    /* Make  $A$  bitonic, sort half forward, half backward */
12    /*  $\otimes$  is exclusive or */
13    sortBitonic( $r \bmod P/2, P/2, n, A, d \otimes (r \geq P/2), \text{False}$ )
14    /* Swap splits keys into lesser and greater halves */
15    /* (or greater and lesser if backward direction) */
16    swapBitonic( $r, P, n, A, d$ )
17    /* Recursive sort of distributed bitonic sequence */
18    sortBitonic( $r \bmod P/2, P/2, n, A, d, \text{True}$ )
```

Algorithm 2: swapBitonic(r, P, n, A, d)

Data: r is the process's rank;
 P is the number of processes (power of 2);
 n is the number of local keys (same on all processes);
 A is the array of keys;
 d is **True** if sorting forward, **False** if sorting backward.

```
/* Communicating partner in the other half */
1  $q \leftarrow r + P/2 \bmod P$ 
2  $A_{\text{recv}} \leftarrow \text{Sendrecv}(A, n, q)$ 
/*  $\otimes$  is exclusive or */
3 if ( $r < q$ )  $\otimes d$  then
    /* Swap to get lesser value */
4     for  $0 \leq i < n$  do
5          $A[i] \leftarrow \min\{A[i], A_{\text{recv}}[i]\}$ 
6 else
    /* Swap to get greater value */
7     for  $0 \leq i < n$  do
8          $A[i] \leftarrow \max\{A[i], A_{\text{recv}}[i]\}$ 
```

2 Runtime analysis

We want an expression for $T_b(n, p)$, the runtime of parallel bitonic sort in terms of n , the number of local keys, and P , the number of processes.

We will first get $\tilde{T}_b(n, p)$, the runtime of parallel sort when the input is bitonic (b is **True**). In this case

$$\tilde{T}_b(n, 1) = T_{lb}(n) := T_{\text{sortLocalBitonic}}(n) \in O(n \log_2 n).$$

Then, the recursive case:

$$\tilde{T}_b(n, P) = (T_{sb}(n, P) := T_{\text{swapBitonic}}(n, P)) + \tilde{T}_b(n, P/2).$$

We will assume that network congestion is not a problem so that $T_{sb}(n, P)$ is independent of P , $T_{sb}(n)$. Ignoring for a moment the cost of local operations in **swapBitonic** (the loops at lines 4 and 7), the cost is at least the cost of a **Sendrecv** of n keys. In the parameters of the LogGP model, the cost of a **Sendrecv** should be

$$T_{sb}(n) \geq L + 2o + G(kn - 1),$$

where L is the network latency, o is the overhead of starting or finishing a communication, G is the inverse of the network bandwidth (i.e., G has units of s/B), and k is the number of bytes per key (8, in our case).

Expanding the recursion, we get

$$\begin{aligned} \tilde{T}_b(n, P) &= \underbrace{T_{sb}(n) + \cdots + T_{sb}(n)}_{\times \log_2 P} + \tilde{T}_b(n, 1) \\ &= T_{sb}(n) \log_2 P + T_{lb}(n). \end{aligned}$$

Now we solve for the full runtime $T_b(n, P)$ when the inputs are not bitonic (b is **False**), the default case.

The base case:

$$T_b(n, 1) = T_l(n) := T_{\text{sortLocal}}(n) \in O(n \log_2 n).$$

The recursive case:

$$\begin{aligned} T_b(n, P) &= T_b(n, P/2) + \tilde{T}_b(n, P) \\ &= \underbrace{\tilde{T}_b(n, P) + \tilde{T}_b(n, P/2) + \cdots + \tilde{T}_b(n, 2)}_{\times \log_2 P} + T_b(n, 1). \end{aligned}$$

We insert our expression for $\tilde{T}_b(n, P)$, group like terms, and use the property of logarithms to turn $\log_2 P/K$ into $\log_2 P - \log_2 K$:

$$\begin{aligned} T_b(n, P) &= T_{sb}(n)(\log_2 P + \log_2 P/2 + \cdots + \log_2 2) + \\ &\quad T_{lb}(n)(\underbrace{1 + 1 + \cdots + 1}_{\times \log_2 P}) + T_b(n, 1) \\ &= T_{sb}(n)(\underbrace{\log_2 P + \log_2 P + \cdots + \log_2 P}_{\times \log_2 P}) - \\ &\quad T_{sb}(n)(0 + 1 + \cdots + (\log_2 P - 1)) + \\ &\quad T_{lb}(n)(\underbrace{1 + 1 + \cdots + 1}_{\times \log_2 P}) + T_b(n, 1) \\ &= T_{sb}(n) \left(\log_2^2 P - \sum_{i=1}^{\log_2 P - 1} i \right) + T_{lb}(n) \log_2 P + T_l(n) \\ &= T_{sb}(n)(\log_2^2 P - \frac{1}{2}(\log_2 P - 1) \log_2 P) + T_{lb}(n) \log_2 P + T_l(n) \\ &= \frac{1}{2} T_{sb}(n)(\log_2^2 P + \log_2 P) + T_{lb}(n) \log_2 P + T_l(n). \end{aligned}$$

We now have a runtime expression for the full **sortBitonic()** function in terms of its components that are simpler to analyze/measure on their own: **sortLocal()**, **sortLocalBitonic()**, and **swapBitonic()**, which we take to be $O(n \log n)$, $O(n \log n)$ and $O(n)$, respectively. If we trust our models of the runtimes of those components, then our model lets us:

- Predict the performance of **sortBitonic()**.
- Verify and validate the performance of **sortBitonic()**. If our measured times do not match the model, the implementation is wrong or the model is wrong (e.g., our pseudocode omits something critical).
- Speculatively compare the performance of other sort implementations to **sortBitonic()**.

If we *don't* already have good models of the component algorithms, we can use the functional form of $T_b(n, P)$ to define an empirical model.

3 Empirical Model

First, let us approximate the components with some as yet undetermined coefficients:

$$\begin{aligned}T_{sb}(n) &\approx C_{sb}^0 + C_{sb}^1 n, \\T_{lb}(n) &\approx C_{lb} n \log_2 n, \\T_l(n) &\approx C_l n \log_2 n.\end{aligned}$$

[We include the constant term C_{sb}^0 because we think network latency may be significant in the case of small n .]

Expanding $T_b(n, P)$:

$$T_b(n, P) \approx \frac{1}{2}(C_{sb}^0 + C_{sb}^1 n)(\log_2^2 P + \log_2 P) + C_{lb} n \log_2 n \log_2 P + C_l n \log_2 n.$$

Now we could measure the runtime of `sortBitonic`(n, P) for a large number of (n, P) pairs, and do least-squares fit for the coefficients, and use a statistical measure of goodness of fit to decide if our four-coefficient model is useful.

We can also isolate some of the coefficient with subsets of the parameter space. If $P = 1$, then C_l (the coefficient for the local sort) is the only one that has an effect, so we can estimate C_l in isolation. If $N = 1$, then we can estimate $(C_{sb}^0 + C_{sb}^1)$ (the minimum cost of performing `swapBitonic()`).